

Special Topics — ai-time-to-die v3 production plan

Evidence status: OBSERVED

Observation cutoff: 2026-08-14T12:00:00-07:00

Topic and current finding

The ai-time-to-die vertical was walked from proposal to complete episodes in one continuous arc (drafts 1-9, KEEP issue #28, episode masters 01-03 published). The v3 production plan converts that walkthrough into a reviewable plan at the same level of detail as the Here Lies Trouble cut-v1 plan: corpus definition, measured episode template, draft-loop history, pilot options, orchestration, gates, and risks. The canonical source corpus is 79.0 seconds, which bounds the series to 2-3 episodes before new material is added.

Evidence and analysis

Measured episode template: 8 segments / 57 s at 1080x1920@24. Whole-episode local render wall time is 3-5 minutes (draft-4..8); draft-9 was a ~0.1-minute audio-only remux because the video was byte-identical to draft-8. Provider spend measured 590 credits total across three Seedance jobs (110 + 240 + 240), with the resolution cap persisting (1080p requested, 720x1280 or 496x864 returned). Voice remap: MeanVC2 with the selected Kokoro female voice zf_xiaoyi (V29, KEEP issue #27) converts the canonical dialogue locally in 1-2 minutes per 30 s window. T5 reuses the here-lies-trouble pilot with original audio; provider-gated episodes cost 110-240 credits each when authorized.

Method, limits, and sources

Method: measured renders over nine drafts, portal review vocabulary (KEEP / KEEP-PARTS / REVISE / REJECT), and the local-object dyad records. Limits: the provider resolution cap persists; PGS subtitle burn remains NOT-BURNED; WASM byte-parity is not yet verified; custom accents require human reference recordings. Sources: izzi repository, situationshipin.space review issues #18-#28, the 29-voice audio bank, the 26-image reference set, and this conversation thread.

PAGE 1 / 4

Coda

Simplified Prompt

Preserves the shortest self-contained assistant response that can be reused to repeat or adapt this special-topic report.

1. [Simplified Prompt](#)

Share

2. [Share summary](#)
3. [Aggregated portable prompt](#)
4. [Sharing boundary](#)
5. [Use AI to change and add your own reassessment info](#)

PAGE 2 / 4

Simplified Prompt

Response clip

Field	Value
Clip ID	ST-CODA-AITTD-V3-001
Source role	assistant
Source reference	izzi Codex thread 2026-08-14 (v3 production plan direction)
Date or date status	2026-08-14
Boundary start	2026-08-14T11:55:00-07:00
Boundary end	2026-08-14T12:05:00-07:00
Clip status	EXACT
Omissions or redactions	NONE
Evidence status	OBSERVED

```
document_type: special-topics
style: randomly-accent
topic: ai-time-to-die vertical production plan v3 (episodes 01-03 + corpus bound)
evidence_boundary: Measured izzi repository records and portal review issues from 2026-08-13/14 only; provider spend bounded to three Seedance jobs (590 credits total).
intended_audience: The project owner and authorized collaborators
output_format: pdf
```

The user approved the v3 plan direction: write proposed_vertical_v3.md at here_lies_trouble_cut_v1 level of detail, synthesize the devastation-pacific-house-style special-topics record (randomly-accent, coda-share) as proposed_vertical_v3.pdf, and publish the plan and PDF for review on situationshipin.space.

PAGE 3 / 4

Share summary

The v3 plan turns the walked-through ai-time-to-die vertical into a reviewable production plan at here_lies_trouble_cut_v1 detail: 8 segments / 57 s template, 79 s canonical corpus bound (2-3 episodes), measured render/voice/provider facts, pilot options, gates, and rizar/ord/eureka orchestration. Production is zero-spend by default; a formal plan-vertical KEEP on situationshipin.space unlocks it.

Aggregated portable prompt

Evidence status: INFERRED — aggregated portable prompt, not transcript.

```
document_type: special-topics
style: randomly-accent
document_type_modifiers:
- coda-share
evidence_scope:
- measured izzi repository records
- portal review issues #18-#28
output_format: pdf
```

Produce ai-time-to-die episodes with the accepted draft-9 pipeline: 8 segments / 57 s / 1080x1920@24, canonical dialogue remapped to zf_xiaoyi via MeanVC2, izzi red-on-white 20 pt two-region bullet text, wild 50% opacity masks, T5 as the HLT stand-in unless provider spend is authorized, shinkansen bed elsewhere, full-motion portal review, validator passing under the 96 MiB budget, estimated against the 79 s canonical corpus bound and the measured 3-5 min per-episode render time.

Sharing boundary

This record may be shared with the project owner and authorized collaborators. Provider prompts and spend receipts remain private and gated.

Use AI to change and add your own reassessment info

Take the aggregated portable prompt from this appendix and combine it with the bundled Devastation Pacific skill file to run your own bounded review.

PAGE 4 / 4